

Department of Electrical Engineering
and Computer Science

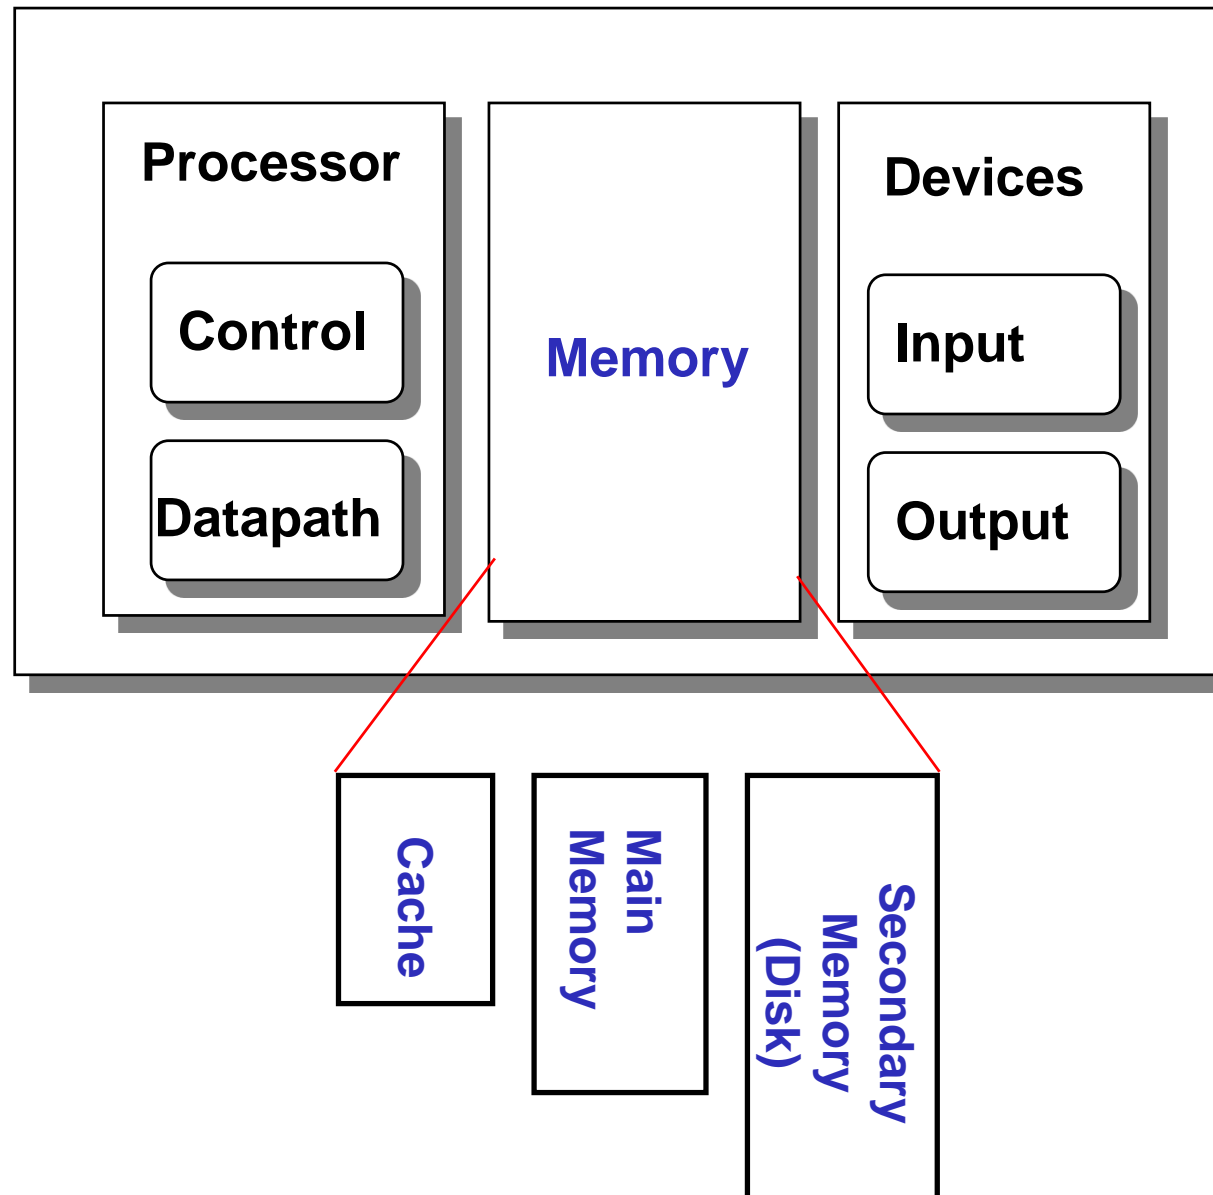


EEL 4768C: Computer Architecture

Memory Hierarchy

Instructor: Suboh A. Suboh

Major Components of a Computer



How is the Hierarchy Managed?


- CPU and RAM are located on the motherboard
- Program Counter register contains address of next instruction to fetch [or for data access, use offset + base from **lw \$rd, offset(base)]**
- CPU outputs bits of PC [or base+offset] over the address bus to RAM and asserts control signals (R/W or enable)
- RAM chip receives address and control, looks up the data at that address, and sends the data bits over the data bus to the CPU
- Animation from 2:20 to 5:30

http://www.youtube.com/watch?v=cNN_tTXABUA&t=2m20sec

How is the Hierarchy Managed?


- registers \leftrightarrow memory
 - by compiler (programmer?)
- cache \leftrightarrow main memory
 - by the cache controller hardware
- main memory \leftrightarrow disks
 - by the operating system (virtual memory)
 - virtual to physical address mapping assisted by the hardware (TLB)
 - by the programmer (files)

Processor-Memory Performance Gap



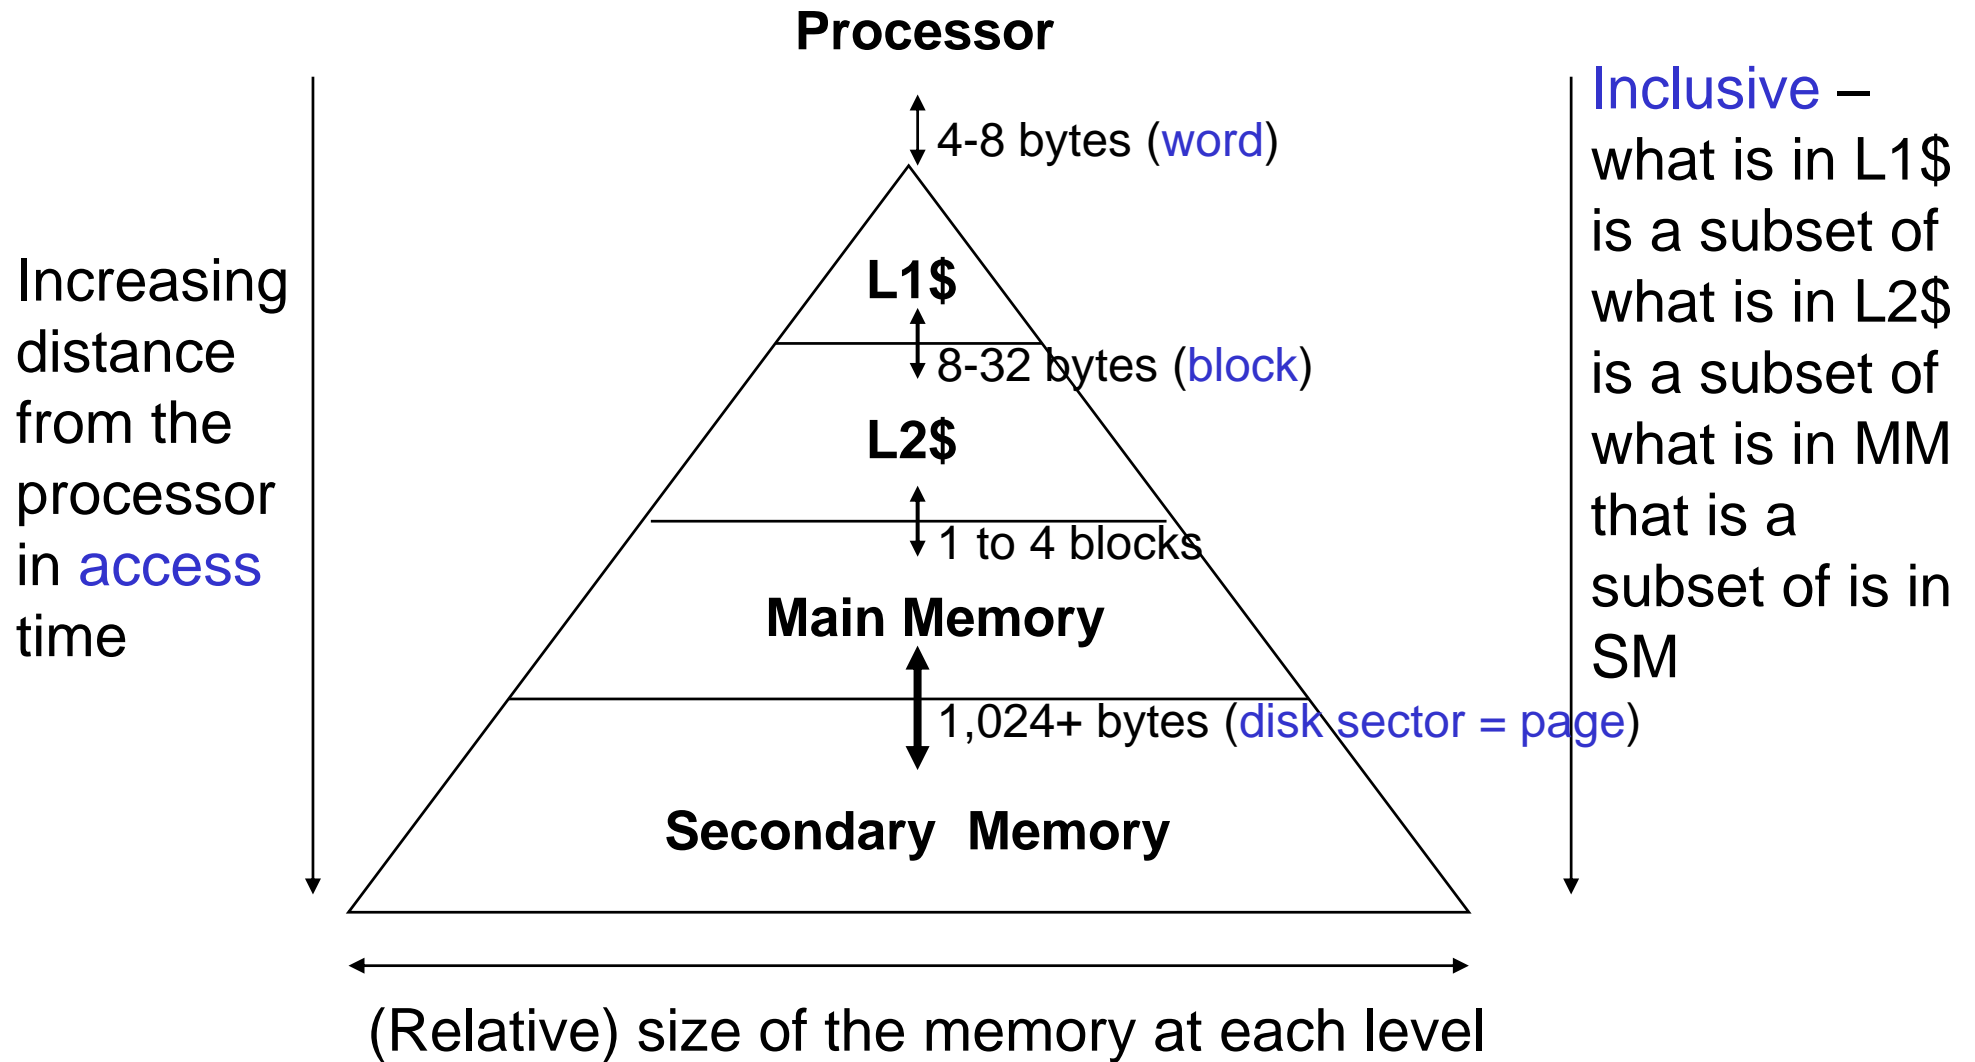
The “Memory Wall”

- Processor vs DRAM speed disparity continues to grow



- Good memory hierarchy (cache) design is increasingly important to overall performance

Memory Hierarchy



The Memory Hierarchy Goal

- Opposing properties of capacity vs speed:
“Large memories are slow and fast memories are small.”
- So how do we create a memory that “gives the **illusion**” of being large, inexpensive, and fast?
 - With **hierarchy** – leverages locality (explained on next page)
 - Except for recent updates, the data at fastest levels is a subset of data at slower levels
 - With **parallelism** – multiple bit access between levels
 - word: 32-bits = 4 bytes transferred at a time (Reg \longleftrightarrow Cache)
 - block: 32 to 128 bytes at a time (Cache \longleftrightarrow Memory)
- We seek to design hardware such that:
 - t_{access} = average access time = close to registers
 - $\$/\text{bit}$ = average cost = close to secondary storage

Principle of Locality

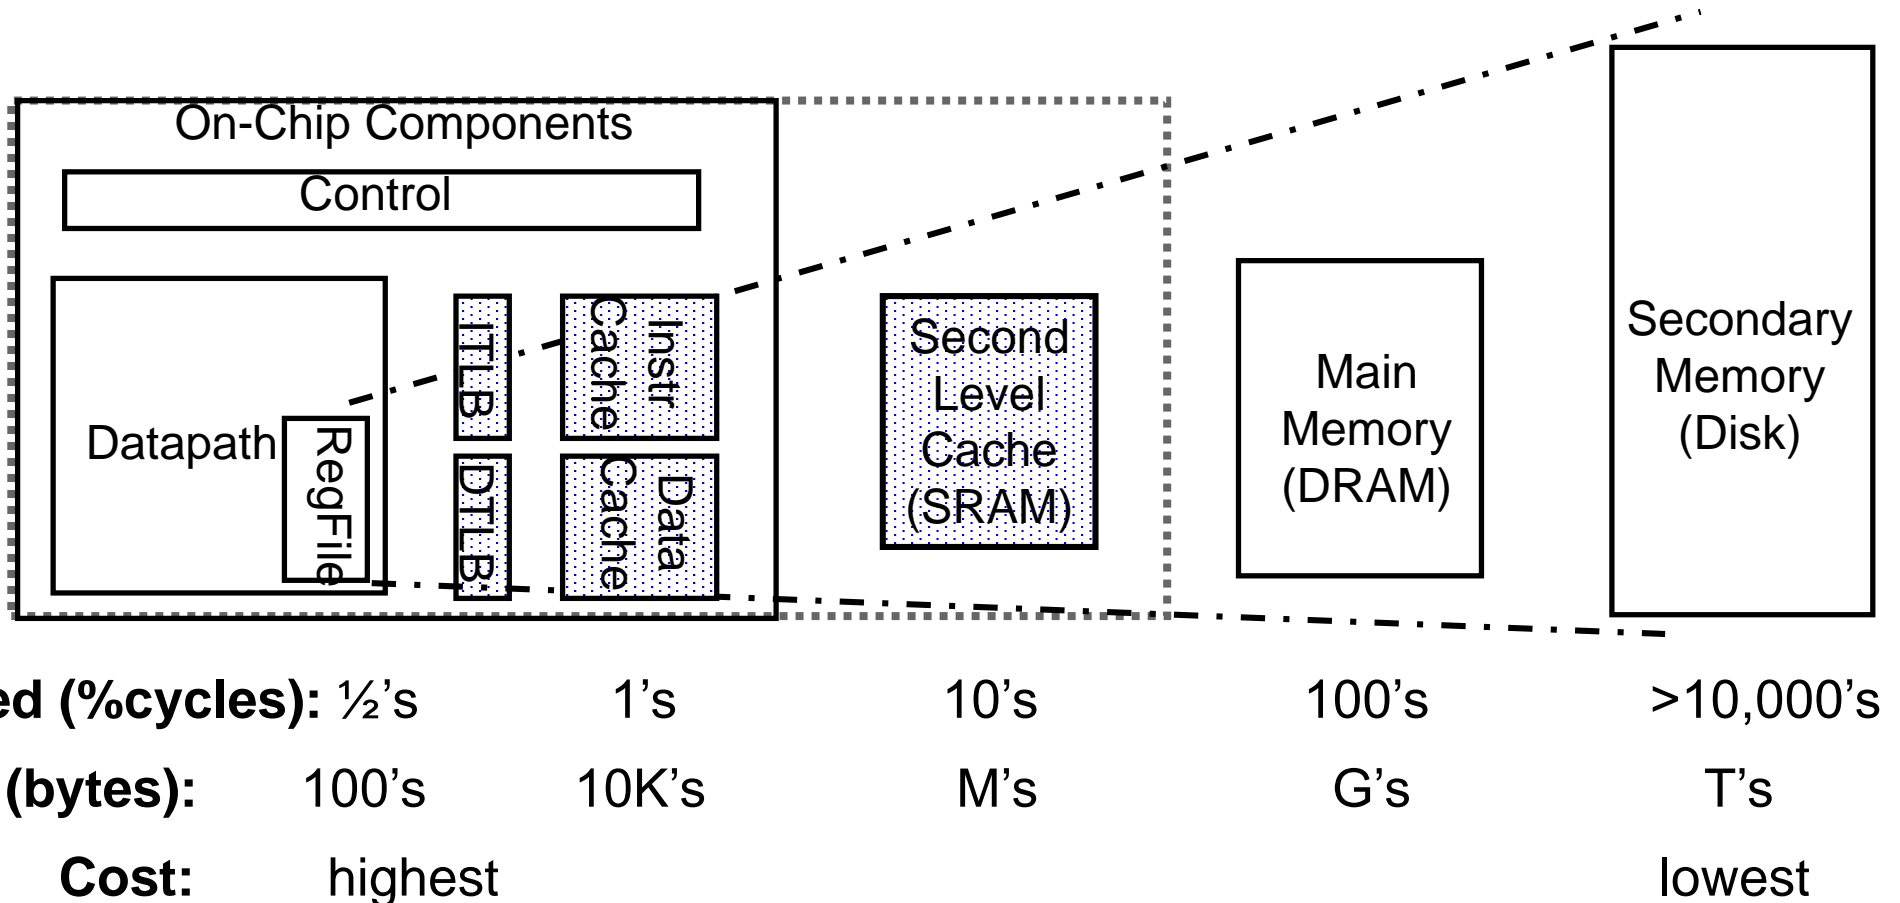
- **Temporal Locality** (locality in time)
 - If a memory location is referenced then it will tend to be referenced again soon
 - ⇒ Keep most recently accessed data items closer to the processor
- **Spatial Locality** (locality in space)
 - If a memory location is referenced, the locations with nearby addresses will tend to be referenced soon
 - ⇒ Move blocks consisting of contiguous words closer to the processor

Taking Advantage of Locality

- Memory hierarchy
- Store everything on disk
- Copy recently accessed (and nearby) items from disk to smaller DRAM memory
 - Main memory
- Copy more recently accessed (and nearby) items from DRAM to smaller SRAM memory
 - Cache memory attached to CPU

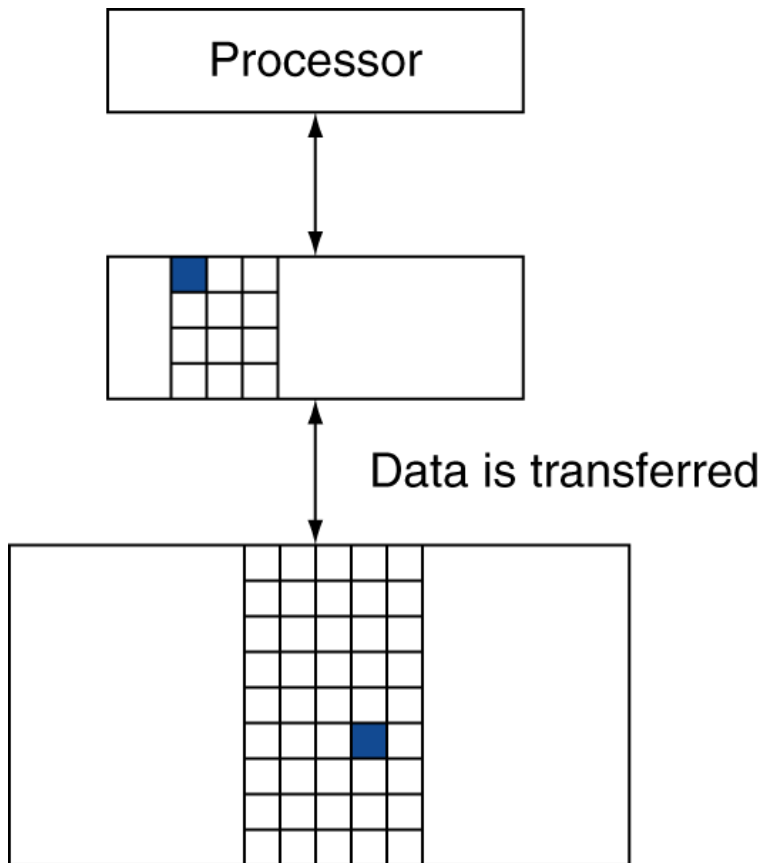
A Typical Memory Hierarchy

Take advantage of the **principle of locality** to present the user with as much memory as is available in the *cheapest* technology at the speed offered by the *fastest* technology



Memory Hierarchy Levels

- Block (aka line): unit of copying
 - May be multiple words
- If accessed data is present in upper level
 - Hit: access satisfied by upper level
 - Hit ratio: hits/accesses
- If accessed data is absent
 - Miss: block copied from lower level
 - Time taken: miss penalty
 - Miss ratio: misses/accesses
 - $= 1 - \text{hit ratio}$
 - Then accessed data supplied from upper level

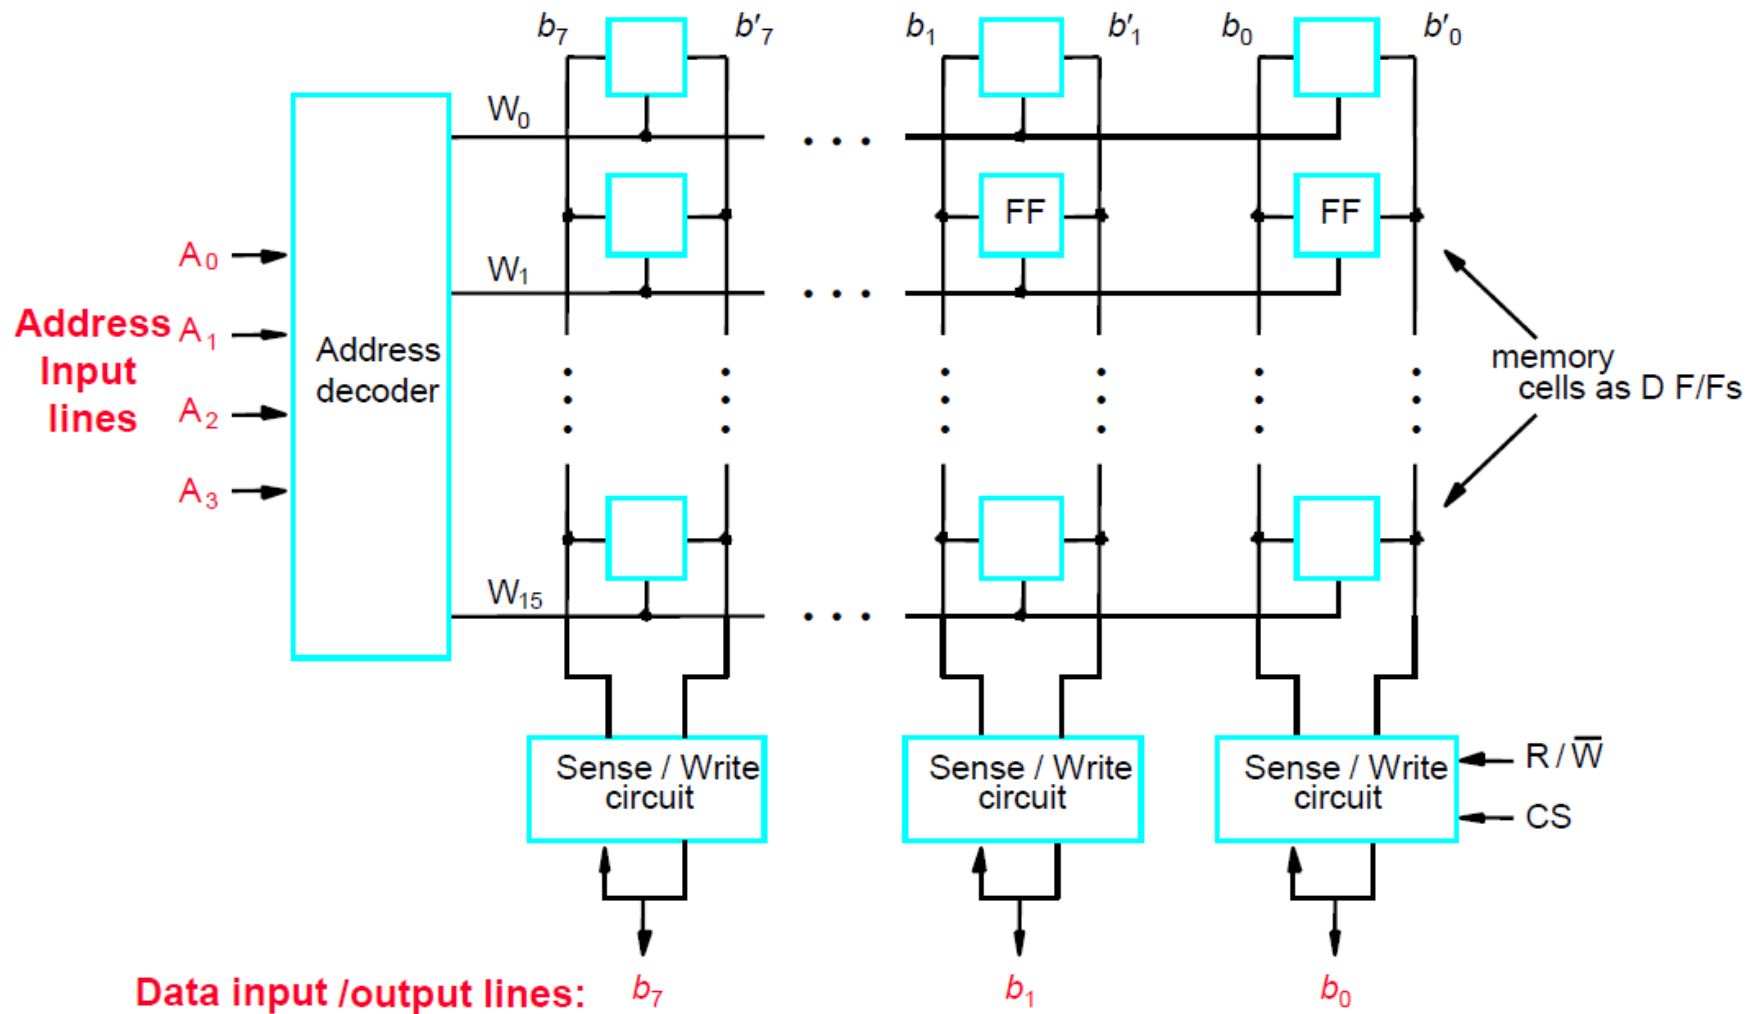


Properties of Memory Devices


Memory Device	Category	Erasure (resetting of all contents to zero)	Write Mechanism (storing a single value)	Volatility (needs power supply to store data)
<i>Random-access memory (RAM)</i>	Read-write memory	Electrically, byte-level	Electrically	Volatile - memories since their contents are lost when power is interrupted
<i>Read-only memory (ROM)</i>	Read-only memory	Not possible	Semiconductor Masks at time of production	Nonvolatile
<i>Programmable ROM (PROM)</i>			Electrically writable	
<i>Erasable PROM (EPROM)</i>	Read-mostly memory	UV light, chip-level	Write time>read time	
<i>Electrically Erasable PROM (EEPROM)</i>		Electrically, byte-level		
<i>Flash memory</i>		Electrically, block-level		

Memory Design – conceptual cell organization


Organization of bit cells in a memory chip.




Increasing Memory Bandwidth



a. One-word-wide memory organization



b. Wider memory organization



c. Interleaved memory organization

1 cycle to transmit the address
15 cycle to access the memory
1 cycle for data transfer

- 4-word wide memory
 - Miss penalty = $1 + 15 + 1 = 17$ bus cycles
 - Bandwidth = $16 \text{ bytes} / 17 \text{ cycles} = 0.94 \text{ B/cycle}$
- 4-bank interleaved memory
 - Miss penalty = $1 + 15 + 4 \times 1 = 20$ bus cycles
 - Bandwidth = $16 \text{ bytes} / 20 \text{ cycles} = 0.8 \text{ B/cycle}$

Cache Memory

- Cache memory
 - The level of the memory hierarchy closest to the CPU
- Given accesses X_1, \dots, X_{n-1}, X_n

X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_3

a. Before the reference to X_n

X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_n
X_3


b. After the reference to X_n

- How do we know if the data is present?
- Where do we look?

The cache just before and just after the reference to a word X_n that is not initially in the cache. This reference caused a miss that forces the cache to fetch X_n from memory and inserted into the cache

Direct Mapped Cache

- Location determined by address
- Direct mapped: only one choice
 - (Block address) modulo (#Blocks in cache)



- #Blocks is a power of 2
- Use low-order address bits

Tags and Valid Bits

- How do we know which particular block is stored in a cache location?
 - Store block address as well as the data
 - Actually, only need the high-order bits
 - Called the *tag*
- What if there is no data in a location?
 - Valid bit: 1 = present, 0 = not present
 - Initially 0

Cache Example

- 8-blocks, 1 word/block, direct mapped
- Initial state

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Miss	110

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
26	11 010	Miss	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Hit	110
26	11 010	Hit	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
16	10 000	Miss	000
3	00 011	Miss	011
16	10 000	Hit	000

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

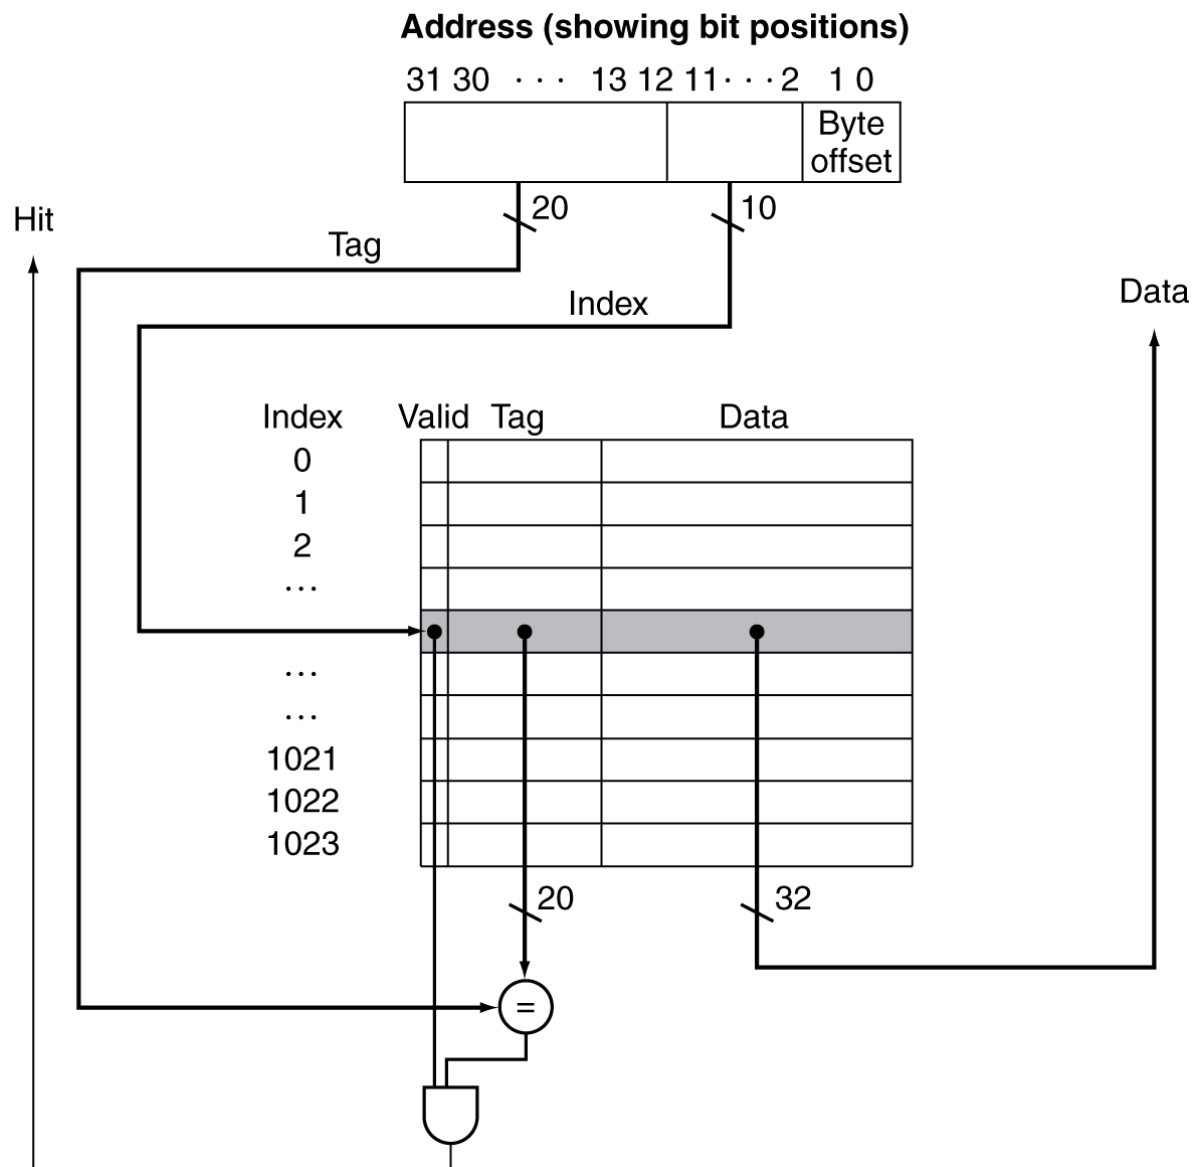
Cache Example

Word addr	Binary addr	Hit/miss	Cache block
18	10 010	Miss	010

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	10	Mem[10010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

Address Subdivision

- 32 bit byte address
- A direct-mapped cache
- The cache size 2^n blocks, n bits used for index
- Block size is 2^m words, (2^{m+2} bytes)



Cache Field Sizes

- The number of bits in a cache includes both the storage for data and for the tags
 - 32-bit byte address
 - For a direct mapped cache with 2^n blocks, n bits are used for the index
 - For a block size of 2^m words (2^{m+2} bytes), m bits are used to address the word within the block and 2 bits are used to address the byte within the word
- *What is the size of the tag field?*
- The total number of bits in a direct-mapped cache is then
$$2^n \times (\text{block size} + \text{tag field size} + \text{valid field size})$$
- *Exercise: How many total bits are required for a direct mapped cache with 16KB of data and 4-word blocks assuming a 32-bit address?*